

Experiment Notes on RLCP

Latest Update: 2024 年 10 月 9 日

摘要

An implementation for localized conformal prediction.

- Author: Rohan Hore, Rina Foygel Barber[HB24].
- RLCP stands for randomly-localized conformal prediction.
- Localized CP: Guan(2023, biometrika)
- Weighted CP: Tibshirani et al. (2019)

1 Simulations

For the following experiments we replicate, our aim will be to achieve 95% coverage, i.e., we choose $\alpha = 0.05$. The feature space is given by $\mathcal{X} = \mathbb{R}^d$. and we can choose a kernel, such as a radial basis kernel.

1.1 Univariate setting

We begin with experiments in a univariate setting, i.e., $d = 1$. We consider the following two data generating distributions.

- Setting 1: $X \sim \mathcal{N}(0, 1), Y \mid X \sim \mathcal{N}\left(\frac{X}{2}, |\sin(X)|\right)$
- Setting 2: $X \sim \mathcal{N}(0, 1), Y \mid X \sim \mathcal{N}\left(\frac{X}{2}, \frac{4}{3}\phi\left(\frac{2X}{3}\right)\right)$, where $\phi(\cdot)$ is the density of a standard Gaussian.

We can use `suppressPackageStartupMessages` to suppress the package startup messages while loading the packages.

```
suppressPackageStartupMessages(library(doParallel))
suppressPackageStartupMessages(library(MASS))
suppressPackageStartupMessages(library(mvtnorm))
suppressPackageStartupMessages(library(ggplot2))
```

The following general code generates the data for the settings.

```

#-----
#-----simulation settings-----
#-----
simulation=function(n,d,setting){
  X=rmvnorm(n,mean=rep(0,d),sigma=diag(d))

  ##setting 1
  if(setting==1){
    Y=0.5*apply(X,1,mean)+apply(abs(sin(X)),1,sum)*rnorm(n,0,1)
  }

  ##setting 2
  if(setting==2){
    Y=0.5*apply(X,1,mean)+apply(2*dnorm(X,0,1.5),1,sum)*rnorm(n,0,1)
  }

  ##setting 3 : P_X uniform on cube
  if(setting==3){
    X=matrix(runif(n*d,-3,3),nrow=n,ncol=d)
    Y=0.5*apply(X,1,mean)+apply(abs(sin(X)),1,sum)*rnorm(n,0,1)
  }

  data=as.data.frame(cbind(Y,X))
  colnames(data)=c("Y",paste0("X",1:d))
  return(data)
}

```

The two settings share the same feature distribution, $P_X = \mathcal{N}(0, 1)$. The difference lies in the conditional distribution of Y , $P_{Y|X}$.

In setting 2, the response has more variance for values of X near the center of the distribution. In contrast, the variance has more variance for values of X lying away from the mean.

```

#-----
#-----Simulating Y|X on a grid of -----
#-----feature points in (-3,3)-----
#-----
conditional_simulation=function(n=100,setting){
  X=as.matrix(seq(-3,3,by=0.01))
  N=rnorm(length(X),0,1)

  ##setting 1
  if(setting==1){Y=0.5*apply(X,1,mean)+abs(sin(X))*N}
  ##setting 2

```

```

if(setting==2){Y=0.5*apply(X,1,mean)+2*dnorm(X,0,1.5)*N}

data=as.data.frame(cbind(Y,X));d=1
colnames(data)=c("Y",paste0("X",1:d))
return(data)
}

```

For both settings, our [score function](#) is given by $s(x, y) = |y - \hat{f}(x)|$, where \hat{f} is a pretrained predictive model. For each setting, we implement the methods using localization kernel H at five different bandwidths, $h \in \{0.1, 0.2, 0.4, 0.8, 1.6\}$; the lowest value represents a highly local H , while the highest value leads to a kernel H that is essentially flat over the bulk of the feature distribution. The methods are run with sample size $n = 2000$, and evaluated on 2000 test points. The entire experiment is repeated for 50 independent trials.

1.2 Generating \tilde{X} for box kernel

Consider a new feature variable \tilde{X}_{n+1} , generated as $\tilde{X}_{n+1} \mid X_{n+1} \sim P_X \circ H(\cdot, X_{n+1})$, where for any function $g : \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$ with $0 < \mathbb{E}_{P_X}[g(X)] < \infty$, we define $P_X \circ g$ as the distribution P_X reweighted by g , i.e.,

$$(P_X \circ g)(A) = \frac{\int_A g(x) dP_X(x)}{\int_{\mathcal{X}} g(x) dP_X(x)} \text{ for all } A \subseteq \mathcal{X}$$

If the localization kernel H has a reasonably [small bandwidth](#), then the reweighted distribution $P_X \circ H(\cdot, X_{n+1})$ [places most of its mass near \$X_{n+1}\$](#) . As a result, we can interpret \tilde{X}_{n+1} as a sort of “synthetic sample” designed to be similar to the test feature X_{n+1} —in a setting where each data point is a patient, we are generating new patient data \tilde{X}_{n+1} for a patient that is [similar](#) (in feature space) to the test patient for whom we want to [perform prediction](#).

The only difference between RLCP and baseLCP, then, is that the [weights](#) are given by the \tilde{w}_i ’s, computed with the kernel H centered at \tilde{X}_{n+1} , rather than the original w_i ’s, which were computed with H centered at X_{n+1} .

$$\begin{aligned} \text{callCP : } w_{i,j} &= \frac{H(X_j, X_i)}{\sum_{k=1}^{n+1} H(X_k, X_i)}, i, j \in [n+1] \\ \text{RLCP : } \tilde{w}_i &= \frac{H(X_i, \tilde{X}_{n+1})}{\sum_{j=1}^{n+1} H(X_j, \tilde{X}_{n+1})}, i \in [n+1] \end{aligned}$$

```

#-----
#-----generating X_tilde for box kernel-----
#-----

runifball=function(n,center,radius){
  d=length(center)
  data=matrix(0,nrow=n,ncol=d)

  U=runif(n,min=0,max=1)

```

```

Z=matrix(rnorm(n*d),nrow=n,ncol=d)
for (i in 1:n){
  data[i,]=center+radius*U[i]^(1/d)*Z[i,]/sqrt(sum(Z[i,]^2))}
return(data)
}

```

2 Base LCP

```

#-----
#-----baseLCP-----
#-----
baseLCP=function(Xcalib,scores_calib,Xtest,scores_test,kernel,h,alpha){
  ntest=dim(Xtest)[1] ## number of test points
  d=dim(Xtest)[2] ## dimension of the feature space
  coverage=score_threshold=rep(0,ntest) ## initializing coverage and score_threshold

  #sorting with respect to the order of calibration scores.
  # scores_calib =  $s(x,y) = |y - \hat{f}(x)|$ 
  Xcalib=as.matrix(Xcalib[order(scores_calib),])
  scores_calib=sort(scores_calib)

  #finding unique scores and the indices where each of these unique scores have been repeated.

  scores=c(scores_calib,Inf)
  indices=list();j=1;i=1
  scores_unique=vector()
  while(i<=length(scores)){
    scores_unique=c(scores_unique,scores[i])
    indices[[j]]=which(scores==scores[i])
    i=i+sum(scores==scores[i]);j=j+1
  }

  for(i in 1:ntest){
    xtest=Xtest[i,];test_score=scores_test[i]
    cov_data=rbind(Xcalib,xtest)

    #finding the weights and the score threshold
    if(kernel=="gaussian"){
      weights=dmvnorm(cov_data,mean=xtest,sigma=diag(d)*h^2)
      result=smoothed_weighted_quantile(scores_unique,alpha,weights,indices)
    }
  }
}

```

```

if(kernel=="box"){
  weights=apply(cov_data,1,FUN=function(x){(euclid_distance(x,xtest)<=h)+0})
  result=smoothed_weighted_quantile(scores_unique,alpha,weights,indices)
}
score_threshold[i]=result[1] #score_threshold
closed=result[2] #whether it's a closed interval

#coverage
coverage[i]=(test_score<score_threshold[i])+0
if(closed==TRUE){coverage[i]=(test_score<=score_threshold[i])+0}
}
return(cbind(coverage,score_threshold))
}

```

This approach compute a quantile $\hat{q}(x)$ that places more weight on the hold out points X_i that lie near x .

- We first choose some *localization kernel* given by a function $H : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$.
- For instance, for $\mathcal{X} = \mathbb{R}^d$, common choices include a Gaussian kernel with bandwidth h given by $H(x, x') = \exp(-\|x - x'\|^2/h^2)$.
- Or a box kernel with bandwidth h given by $H(x, x') = 1\{\|x - x'\| \leq h\}$.
- The local quantile is then computed as

$$\hat{q}_{1-\alpha}(X_{n+1}) = \text{Quantile}_{1-\alpha} \left(\sum_{i=1}^n w_i \delta_{s(X_i, Y_i)} + w_{n+1} \delta_{\infty} \right).$$

- δ_s is the point mass at s .
- weights are given by

$$w_i = \frac{H(X_i, X_{n+1})}{\sum_{j=1}^{n+1} H(H_j, X_{n+1})}, \quad i \in [n+1].$$

- The resulting prediction interval is defined as

$$\hat{C}_n^{\text{baseLCP}}(X_{n+1}) = \{y \in \mathcal{Y} : s(X_{n+1}, y) \leq \hat{q}_{1-\alpha}(X_{n+1})\}$$

- We choose the trivial localization kernel $H(x, x') = 1$, namely, no localization, then we would have $\hat{q}_{1-\alpha}(x)$ constant over x .

2.1 Smooted Weighted Quantile

```

#-----
#---computing smoothed weighted quantile-----
#-----
##---here v consists of the unique scores, while indices contain the data-indices that

```

```

##----has led to same score, w is the weight vector.
smoothed_weighted_quantile=function(v,alpha,w,indices){
  w=w/sum(w) ## normalizing the weights by definition
  U=runif(1,min=0,max=1)

  #finding the weights corresponding to each unique score.
  v_tilde=v ## unique scores
  w_tilde=rep(0,length(v))
  for(i in 1:length(v)){
    w_tilde[i]=sum(w[indices[[i]]]) ## summing the weights corresponding to each unique score.
  }

  #computing p-value at points in between the calibration scores.
  p_values=rep(0,length(v_tilde))
  for(i in 1:length(v_tilde)-1){
    p_values[i]=sum(w_tilde[i:(length(v_tilde)-1)])+U*(tail(w_tilde,1))
  }
  #computing p-value at a point higher than all calibration scores.
  p_values[length(v_tilde)]=U*(tail(w_tilde,1))

  #if pvalue is never greater than alpha, we output the empty set.
  if(sum((p_values>alpha))>0){
    id=max(which(p_values>alpha))
    #now we check, whether the prediction interval will be a closed interval or open.
    quantile=v_tilde[id]
    if(id<length(v_tilde)-1){closed=(sum(w_tilde[(id+1):(length(v_tilde)-1)])+U*(w_tilde[id]+ta
    if(id==length(v_tilde)){closed=FALSE}
    if(id==length(v_tilde)-1){closed=(U*(w_tilde[id]+tail(w_tilde,1))>alpha)}}
  else{quantile=-Inf;closed=FALSE}
  return(c(quantile,closed))
}

```

References

- [HB24] Rohan Hore and Rina Foygel Barber. *Conformal prediction with local weights: randomization enables local guarantees*. 2024. arXiv: 2310.07850 [stat.ME]. URL: <https://arxiv.org/abs/2310.07850>.